Quantifying the AI Access Gap

One-page findings brief · Updated evidence cutoff: October 13, 2026

What can be defended

The gap depends strongly on the comparison population. No universal physical-compute multiplier is identified.

Benchmark	Result	Essential qualification
OpenAI research expenditure	P50 ~$601/day; P90 ~$7,047/day	Retail-equivalent weekly means of daily quantiles; broad research staff, not all inference
Selected public coding users, June	P50 $29/coding day; P90 $215; P99 $708	Viberank submitted logs; highest-cost windows selected
Internal / selected-user expenditure	P50 benchmark **~21x**; P90 **~33x**	Different dates, activity, selection and price bases
Other external developer benchmarks	Claude enterprise mean $13/active day; Jellyfish median $52.38/month	These yield very different comparisons; neither is a matched daily median
Codex same-product runtime	Corresponding-percentile gaps **~4.5-7.1x**	Active accounts; all internal employees; runtime includes tools and short waits

Sources: OpenAI research, Viberank June, Claude costs, Jellyfish, Codex paper.

New measured token anchors

OpenAI's sampled, capped estimate is **~723,000 output tokens/research employee/day**. Viberank's matching rounded totals imply **~202,000 output tokens/active coding day**, giving a **~3.6x heterogeneous benchmark**. A small academic study implies **~1,005 output tokens/student/calendar day** for 37 selected users on one platform, giving **~719x**. These distinct scopes do not identify a casual-user median or a physical-compute ratio.

The October Viberank snapshot reports a selected-window median of **$1,521 per normalized month**. Dividing by 30 yields a daily rate, not a daily quantile. Its lab affiliations are unverified; it is not a representative external-professional sample.

What remains missing

The broad consumer study provides **~3.67 messages/weekly-active account/calendar day**, without full tokens or a casual median. Matched model, cache, hidden-reasoning and person-day records remain absent across tiers. No reliable FLOP, GPU-hour or energy multiplier can be published. Karpathy's full original-versus-edited long-post wording remains only partially primary-verified.

Article-ready conclusion

The evidence supports substantial, uneven use of AI. It supports selected expenditure benchmarks from tens to hundreds of times, alongside smaller same-product runtime and some output-token comparisons. Those quotients describe different populations and accounting rules; they are not bounds on one universal gap. **Do not substitute subscriptions for consumption, agent-hours for GPU-hours, or greater token volume for proportionately greater capability.**

The expanded package contains 45 source records, 5,353 observations and eight figure pairs. Missing values and conflicting Copilot totals remain explicit; scenario ranges are not confidence intervals.